




Data Boot Camp
Lesson 4.2



Class Objectives

By the end of today's class, you will be able to:



Navigate through DataFrames using `loc` and `iloc`.



Filter and slice Pandas DataFrames.



Create and access Pandas `groupby` objects.




Sort DataFrames.



Instructor Demonstration

Exploring Data with loc and iloc



Programmers can easily collect specific rows and columns of data from a DataFrame by using the `loc` and `iloc` methods.

Exploring Data with `loc` and `iloc`

- `loc` returns data based on an index of labels/strings
- `loc` is limited to string types and cannot be used on a numerical index. As an alternative solution, you can use the `df.set_index()` function, passing in the desired column header for the index.
- Instead of using labels, `iloc` uses integer-based indexing for selection by position.

```
# Set new index to STREET NAME
df = original_df.set_index("STREET NAME")
df.head()
```

	STREET NAME ID	STREET FULL NAME	POSTAL COMMUNITY	MUNICIPAL COMMUNITY
STREET NAME				
PRIVATE STREET	1400342	PRIVATE STREET	BATON ROUGE	BATON ROUGE
4TH	1	N 4TH ST	BATON ROUGE	BATON ROUGE
11TH	10	S 11TH ST	BATON ROUGE	BATON ROUGE
ADDINGTON	100	ADDINGTON AVE	BATON ROUGE	BATON ROUGE
CHALFONT	1000	W CHALFONT DR	BATON ROUGE	PARISH


Exploring Data with `loc` and `iloc`

- Both `loc` and `iloc` use brackets that contain the desired rows, followed by a comma and the desired columns.
- For example: `loc["ADDINGTON", "STREET FULL NAME"]`

gives the same result as: `iloc[3,1]`

```
# Set new index to STREET NAME
df = original_df.set_index("STREET NAME")
df.head()
```

	STREET NAME ID	STREET FULL NAME	POSTAL COMMUNITY	MUNICIPAL COMMUNITY
STREET NAME				
PRIVATE STREET	1400342	PRIVATE STREET	BATON ROUGE	BATON ROUGE
4TH	1	N 4TH ST	BATON ROUGE	BATON ROUGE
11TH	10	S 11TH ST	BATON ROUGE	BATON ROUGE
ADDINGTON	100	ADDINGTON AVE	BATON ROUGE	BATON ROUGE
CHALFONT	1000	W CHALFONT DR	BATON ROUGE	PARISH



`loc` and `iloc` can be used to conditionally filter rows of data based on the values within a column.



Activity: Good Movies

In this activity, you will create an application that searches through IMDb data to find only the best movies out there.

Suggested Time:

20 minutes




Time's Up! Let's Review.



Instructor Demonstration

Cleaning Data



**When dealing with massive datasets,
it's almost inevitable that we'll
encounter duplicate rows, inconsistent
spelling, and missing values.**

Cleaning Data

```
del <DataFrame>[<columns>]
```

```
In [4]: # Preview of the DataFrame
# Note that Memo_CD is likely a meaningless column
df.head()
```

Out[4]:

	Name	Employer	City	State	Zip	Amount	Memo_CD
0	CAREY, JAMES	NOT EMPLOYED	HOCKESSIN	DE	197071618.0	500	NaN
1	OBICI, SILVANA	STONY BROOK	PORT JEFFERSON STATION	NY	117764286.0	250	NaN
2	MAISLIN, KAREN	RETIRED	WILLIAMSVILLE	NY	14221.0	250	NaN
3	MCCLELLAND, CARTER AND STEPHANIE	UNION SQUARE ADVISORS	NEW YORK	NY	10023.0	1000	NaN
4	MCCLUSKEY, MARTHA	STATE UNIVERSITY OF NEW YORK	BUFFALO	NY	14214.0	250	NaN

```
In [5]: # Delete extraneous column
del df['Memo_CD']
df.head()
```

Out[5]:

	Name	Employer	City	State	Zip	Amount
0	CAREY, JAMES	NOT EMPLOYED	HOCKESSIN	DE	197071618.0	500
1	OBICI, SILVANA	STONY BROOK	PORT JEFFERSON STATION	NY	117764286.0	250
2	MAISLIN, KAREN	RETIRED	WILLIAMSVILLE	NY	14221.0	250
3	MCCLELLAND, CARTER AND STEPHANIE	UNION SQUARE ADVISORS	NEW YORK	NY	10023.0	1000
4	MCCLUSKEY, MARTHA	STATE UNIVERSITY OF NEW YORK	BUFFALO	NY	14214.0	250

Cleaning Data

count()

<DataFrame>.dropna(how='any')

```
In [6]: # Identify incomplete rows
df.count()
```

```
Out[6]: Name      2000
Employer    1820
City        1999
State       1999
Zip         1996
Amount      2000
dtype: int64
```

```
In [7]: # Drop all rows with missing information
df = df.dropna(how='any')
```

```
In [8]: # Verify dropped rows
df.count()
```

```
Out[8]: Name      1818
Employer    1818
City        1818
State       1818
Zip         1818
Amount      1818
dtype: int64
```

Cleaning Data

value_counts()

replace()

```
In [12]: # Display an overview of the Employers column
df['Employer'].value_counts()
```

```
Out[12]: NOT EMPLOYED          609
         NONE                 321
         SELF-EMPLOYED        132
         SELF                  33
         RETIRED              32
         ...
         INTEL CORPORATION      1
         SLOCUM & SONS          1
         OCPS                   1
         HEALTHCARE PARTNERS    1
         CARBON FIVE           1
         Name: Employer, Length: 519, dtype: int64
```

```
In [13]: # Clean up Employer category. Replace 'SELF' and 'SELF EMPLOYED' with 'SELF-EMPLOYED'
df['Employer'] = df['Employer'].replace({'SELF': 'SELF-EMPLOYED', 'SELF EMPLOYED': 'SELF-EMPLOYED'})
```

```
In [14]: # Verify clean-up.
df['Employer'].value_counts()
```

```
Out[14]: NOT EMPLOYED          609
         NONE                 321
         SELF-EMPLOYED        180
         RETIRED              32
         INGRAM BARGE COMPANY  30
```



Activity: Hong Kong LPG Appliances

In this activity, you will take an LPG appliance dataset from Hong Kong, and clean it so that the DataFrame is consistent, and that there are no rows with missing data.

Suggested Time:

20 minutes



Time's Up! Let's Review.



Activity: Pandas Recap and Data Types

In this activity, we will recap what has been covered in Pandas up to this point.

Suggested Time:

15 minutes



A close-up photograph of a white computer keyboard. The central focus is a single white key with rounded edges. On this key, there is a dark blue icon of a coffee cup with three wavy lines above it representing steam. Below the icon, the word "Break" is printed in a dark blue, serif font. The key is set against a light-colored, textured keyboard surface. Other keys are visible in the background, including one with a double quote symbol and another with a tilde symbol, but they are out of focus.

Break



Instructor Demonstration

Pandas Grouping



`groupby` is a simpler
method for filtering data.

Pandas Grouping

To split the DataFrame into multiple groups and group by state, we use `df.groupby([<Columns>])`.

The `groupby` method returns a `groupby` object that can only be accessed by using a data function on it.

```
# Count how many loss incidents occurred in each city
grouped_city_df = loss_df.groupby(["Incident City"])

print(grouped_city_df)

grouped_city_df.count().head(10)
```

<pandas.core.groupby.generic.DataFrameGroupBy object at 0x7fd919ddadf0>

	Fire Department Name	Incident date	Incident Type Code	Incident Type	Alarm Date and Time	Arrival Date and Time	Last Unit Cleared Date and Time	Property Loss	Contents Loss	Fire Service Deaths	Fire Service Injuries	Other Fire Deaths	Other Fire Injuries	Incident Zip Code	Response Time (seconds)	In De
Incident City																
AMSTON	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
ANSONIA	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2
AVON	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6
Andover	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3
Ansonia	18	18	18	18	18	18	18	18	18	18	18	18	18	18	18	18
BERLIN	32	32	32	32	32	32	32	32	32	32	32	32	32	32	32	32
BETHEL	13	13	13	13	13	13	13	13	13	13	13	13	13	13	13	13
BLOOMFIELD	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8
BRANFORD	21	21	21	21	21	21	21	21	21	21	21	21	21	21	21	21
BRIDGEPORT	13	13	13	13	13	13	13	13	13	13	13	13	13	13	13	13

Pandas Grouping

The `pd.DataFrame()` method makes it possible to create new DataFrames by using only `groupby` data.

A DataFrame can also be created by selecting a single Series from a `groupby` object and passing it in as the values for a specified column.

```
# Save loss sums as series
city_property_loss = grouped_city_df["Property Loss"].sum()
city_contents_loss = grouped_city_df["Contents Loss"].sum()
city_contents_loss.head()
```

```
Incident City
AMSTON      5000.0
ANSONIA      600.0
AVON       1250.0
Andover      500.0
Ansonia    265100.0
Name: Contents Loss, dtype: float64
```

```
# Create a new DataFrame using count and loss amounts
city_summary_df = pd.DataFrame({"Number of Loss Incidents": city_loss_counts,
                                "Total Property Loss": city_property_loss,
                                "Total Contents Loss": city_contents_loss})
city_summary_df.head()
```

	Number of Loss Incidents	Total Property Loss	Total Contents Loss
AMSTON	1	65000.0	5000.0
ANSONIA	2	5000.0	600.0
AVON	6	14200.0	1250.0
Andover	3	2500.0	500.0
Ansonia	18	644100.0	265100.0

Pandas Grouping



It's also possible to perform a `df.groupby()` method on multiple columns by passing two or more column references into the list parameter.

```
# It is also possible to group a DataFrame by multiple columns
# This returns an object with multiple indexes, however, which can be harder to deal with
grouped_city_loss_incidents = loss_df.groupby(["Incident City", "Incident Type Code"])
grouped_city_loss_incidents.count().head(10)
```

		Fire Department Name	Incident date	Incident Type	Alarm Date and Time	Arrival Date and Time	Last Unit Cleared Date and Time	Property Loss	Contents Loss	Fire Service Deaths	Fire Service Injuries	Other Fire Deaths	Other Fire Injuries
Incident City	Incident Type Code												
AMSTON	111	1	1	1	1	1	1	1	1	1	1	1	1
ANSONIA	111	2	2	2	2	2	2	2	2	2	2	2	2
AVON	111	3	3	3	3	3	3	3	3	3	3	3	3
	113	1	1	1	1	1	1	1	1	1	1	1	1
	114	1	1	1	1	1	1	1	1	1	1	1	1
	131	1	1	1	1	1	1	1	1	1	1	1	1
Andover	111	2	2	2	2	2	2	2	2	2	2	2	2
	113	1	1	1	1	1	1	1	1	1	1	1	1
Ansonia	111	12	12	12	12	12	12	12	12	12	12	12	12
	113	2	2	2	2	2	2	2	2	2	2	2	2

Pandas Grouping



A new DataFrame can be created from a `groupby` object.

```
# Converting a GroupBy object into a DataFrame
total_city_loss_df = pd.DataFrame(
    grouped_city_loss_incidents[["Property Loss", "Contents Loss"]].sum())
total_city_loss_df.head(10)
```

		Property Loss	Contents Loss
Incident City	Incident Type Code		
AMSTON	111	65000.0	5000.0
ANSONIA	111	5000.0	600.0
AVON	111	8000.0	1200.0
	113	0.0	50.0
	114	1000.0	0.0
	131	5200.0	0.0
Andover	111	2500.0	200.0
	113	0.0	300.0
Ansonia	111	617100.0	263500.0
	113	5000.0	500.0



Activity: Exploring U.S. Census Data

In this activity, you will revisit the U.S. Census data and create DataFrames with calculated totals and averages of each state by year.

Suggested Time:

25 minutes



Time's Up! Let's Review.



Instructor Demonstration

Sorting Made Easy

Sorting with Pandas



To sort a DataFrame based on the values within a column, use the `df.sort_values()` method and pass in the column name to sort by as a parameter.



The "ascending" parameter is always marked as True by default. Therefore, the `sort_values()` method will always sort from lowest to highest unless the parameter of `ascending=False` is also passed into the `sort_values()` method.



Activity: Search for the Worst

In this activity, you will take a dataset on San Francisco Airports' utility consumption and determine which day in the dataset had the worst consumption for each utility.

Suggested Time:

20 minutes



Time's Up! Let's Review.

Questions?

